

457final.R

hyukjang

2025-03-25

```
# Load required libraries
library(readr)      # For reading CSV files
library(lubridate)   # For date handling

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

library(forecast)    # For ARIMA modeling and forecasting

## Warning: package 'forecast' was built under R version 4.3.3

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

library(ggplot2)      # For plotting
library(tseries)      # For stationarity tests

## Warning: package 'tseries' was built under R version 4.3.3

# 1. Data Import and Preprocessing

# Read the CSV file (ensure the working directory is set correctly or provide full path)
data <- read_csv("Daily Prices_ICCO (1).csv")

## Rows: 7812 Columns: 2

## -- Column specification -----
## Delimiter: ","
## chr (1): Date
## num (1): ICCO daily price (US$/tonne)
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

# Examine structure (adjust column names if necessary)
str(data)

## spc_tbl_ [7,812 x 2] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
##   $ Date                : chr [1:7812] "27/02/2025" "26/02/2025" "25/02/2025" "24/02/2025" ..
##   $ ICCO daily price (US$/tonne): num [1:7812] 9100 9090 8669 8409 9106 ...
##   - attr(*, "spec")=
##     .. cols(
```

```
## .. Date = col_character(),
## .. `ICCO daily price (US$/tonne)` = col_number()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
head(data)
```

```
## # A tibble: 6 x 2
##   Date           `ICCO daily price (US$/tonne)`
##   <chr>                                <dbl>
## 1 27/02/2025                9100.
## 2 26/02/2025                9090.
## 3 25/02/2025                8669.
## 4 24/02/2025                8409.
## 5 21/02/2025                9106.
## 6 20/02/2025                9962.
```

```
# Create a time series object.
```

```
# For daily data, set frequency = 365 (or use an appropriate frequency if weekends/holidays are removed.
```

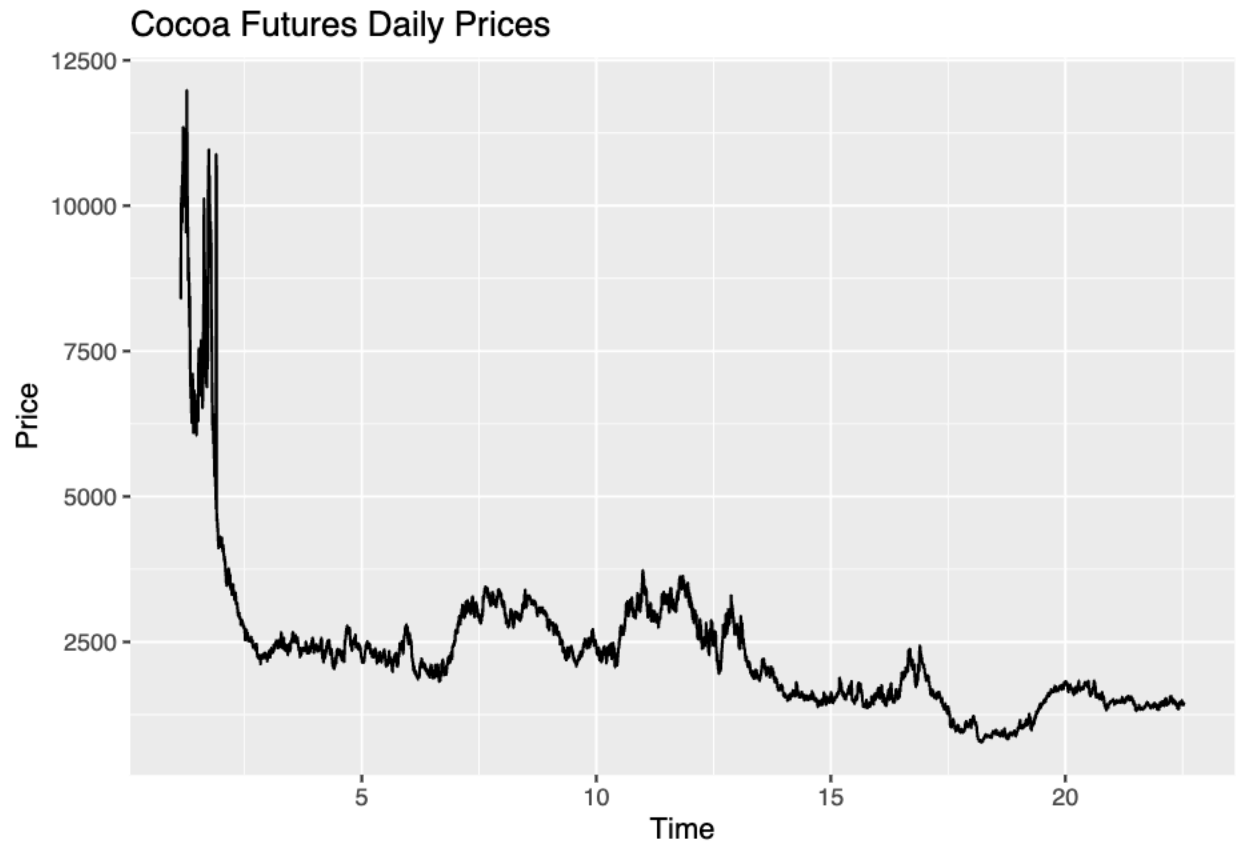
```
start_year <- year(min(data$Date))
```

```
start_doy <- yday(min(data$Date))
```

```
price_ts <- ts(data$`ICCO daily price (US$/tonne)`, frequency = 365, start = c(start_year, start_doy))
```

```
# Plot the raw time series
```

```
autoplot(price_ts) +
  ggtitle("Cocoa Futures Daily Prices") +
  xlab("Time") +
  ylab("Price")
```



2. Stationarity Check & Differencing

Perform the Augmented Dickey-Fuller Test to check for stationarity

```
adf_result <- adf.test(price_ts)
```

```
## Warning in adf.test(price_ts): p-value smaller than printed p-value
```

```
print(adf_result)
```

```
##
```

```
## Augmented Dickey-Fuller Test
```

```
##
```

```
## data: price_ts
```

```
## Dickey-Fuller = -6.6276, Lag order = 19, p-value = 0.01
```

```
## alternative hypothesis: stationary
```

If the p-value is > 0.05, the series is likely non-stationary so we difference it

```
price_diff <- diff(price_ts)
```


```
autoplot(price_diff) +
```

```
  ggtitle("First Difference of Cocoa Futures Prices") +
```

```
  xlab("Time") +
```


```
  ylab("Differenced Price")
```

First Difference of Cocoa Futures Prices




```
# Plot ACF and PACF for the differenced series (use these plots to help decide on the ARIMA orders)  
acf(price_diff, main="ACF of Differenced Prices")
```

ACF of Differenced Prices



```
pacf(price_diff, main="PACF of Differenced Prices")
```

PACF of Differenced Prices



3. Candidate ARIMA Models

*# Based on your ACF/PACF interpretation (e.g., a spike at lag 1 in both ACF and PACF),
try fitting ARIMA(1,1,0), ARIMA(0,1,1), and ARIMA(1,1,1).*

```
arima110 <- arima(price_ts, order = c(1, 1, 0))
```

```
arima011 <- arima(price_ts, order = c(0, 1, 1))
```

```
arima111 <- arima(price_ts, order = c(1, 1, 1))
```

Summaries of the models to inspect parameter estimates and diagnostic metrics

```
summary(arima110)
```

```
##
```

```
## Call:
```

```
## arima(x = price_ts, order = c(1, 1, 0))
```

```
##
```

```
## Coefficients:
```

```
##          ar1
```

```
##        -0.5790
```

```
## s.e.    0.0092
```

```
##
```

```
## sigma^2 estimated as 15857:  log likelihood = -48854.98,  aic = 97713.97
```

```
##
```

```
## Training set error measures:
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
```

```
## Training set -1.541816 125.9154 39.69301 -0.07304487 1.456403 1.146831
```

```
##
```

```
##          ACF1
```

```
## Training set 0.04704558
```

```
summary(arima011)
```

```
##
## Call:
## arima(x = price_ts, order = c(0, 1, 1))
##
## Coefficients:
##          ma1
##        -0.4575
## s.e.    0.0087
##
## sigma^2 estimated as 17977:  log likelihood = -49344.95,  aic = 98693.89
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -1.799099 134.0682 39.65113 -0.08647213 1.433988 1.145621
##              ACF1
## Training set -0.1287615
```

```
summary(arima111)
```

```
##
## Call:
## arima(x = price_ts, order = c(1, 1, 1))
##
## Coefficients:
##          ar1      ma1
##        -0.6476  0.1026
## s.e.    0.0132  0.0161
##
## sigma^2 estimated as 15778:  log likelihood = -48835.59,  aic = 97677.19
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -1.458705 125.6032 39.13152 -0.06952958 1.434941 1.130609
##              ACF1
## Training set 0.007406219
```

4. Model Comparison Using AIC

```
# Create a table with AIC values for comparison
```

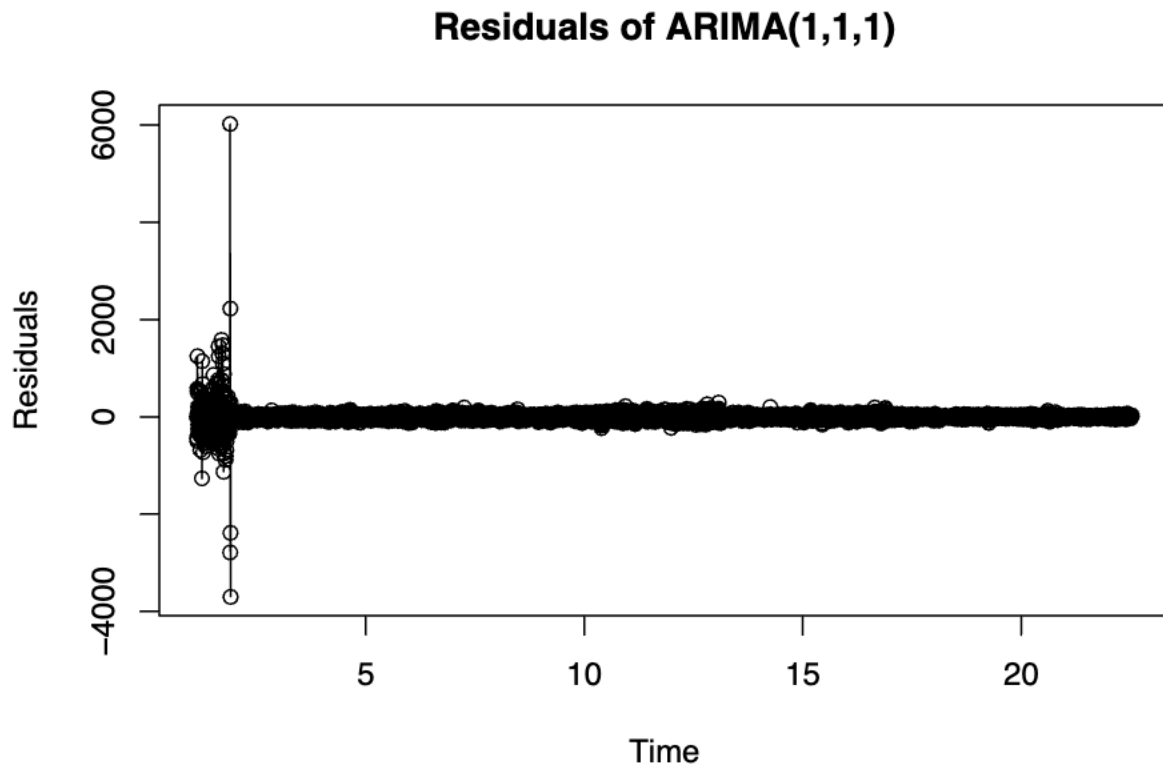
```
model_comparison <- data.frame(
  Model = c("ARIMA(1,1,0)", "ARIMA(0,1,1)", "ARIMA(1,1,1)"),
  AIC = c(AIC(arima110), AIC(arima011), AIC(arima111))
)
print(model_comparison)
```

```
##           Model      AIC
## 1 ARIMA(1,1,0) 97713.97
## 2 ARIMA(0,1,1) 98693.89
## 3 ARIMA(1,1,1) 97677.19
```

5. Residual Diagnostics for the Chosen Model


```
# Suppose the AIC values suggest that ARIMA(1,1,1) is best; then check its residuals:
```

```
# Residual plot over time  
plot(residuals(arima111), type = "o",  
     main = "Residuals of ARIMA(1,1,1)",  
     xlab = "Time", ylab = "Residuals")
```




```
# ACF of residuals to check for autocorrelation  
acf(residuals(arima111), main="ACF of ARIMA(1,1,1) Residuals")
```


ACF of ARIMA(1,1,1) Residuals



```
# Q-Q plot for normality of residuals  
qqnorm(residuals(arima111))  
qqline(residuals(arima111))
```

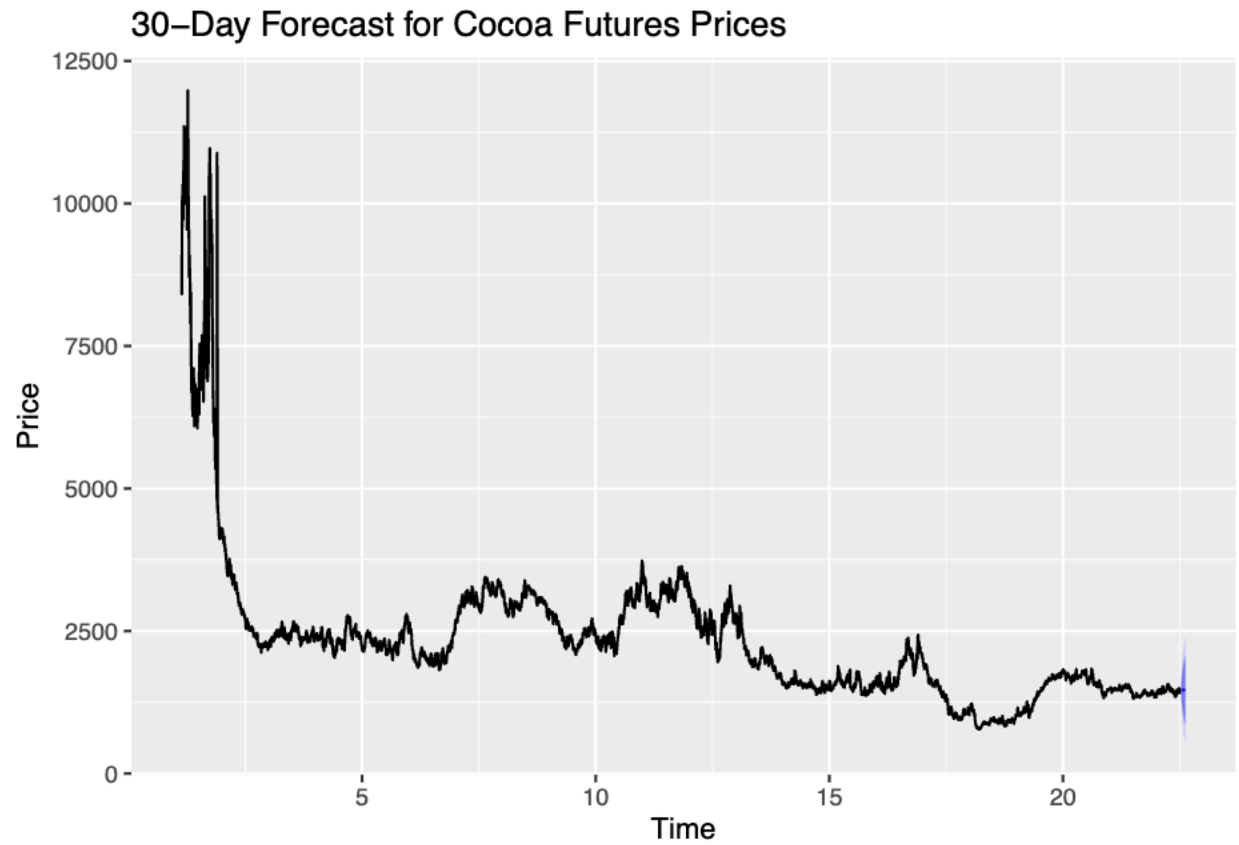
Normal Q-Q Plot



```
# Ljung-Box test to statistically test for autocorrelation in residuals
lb_test <- Box.test(residuals(arima111), lag = 10, type = "Ljung-Box")
print(lb_test)

##
## Box-Ljung test
##
## data: residuals(arima111)
## X-squared = 300.18, df = 10, p-value < 2.2e-16

# 6. Forecasting
# Forecast the next 30 days (or adjust as needed) using the selected model
forecast_horizon <- 30
fc <- forecast(arima111, h = forecast_horizon)
autoplot(fc) +
  ggtitle("30-Day Forecast for Cocoa Futures Prices") +
  xlab("Time") +
  ylab("Price")
```



End of script